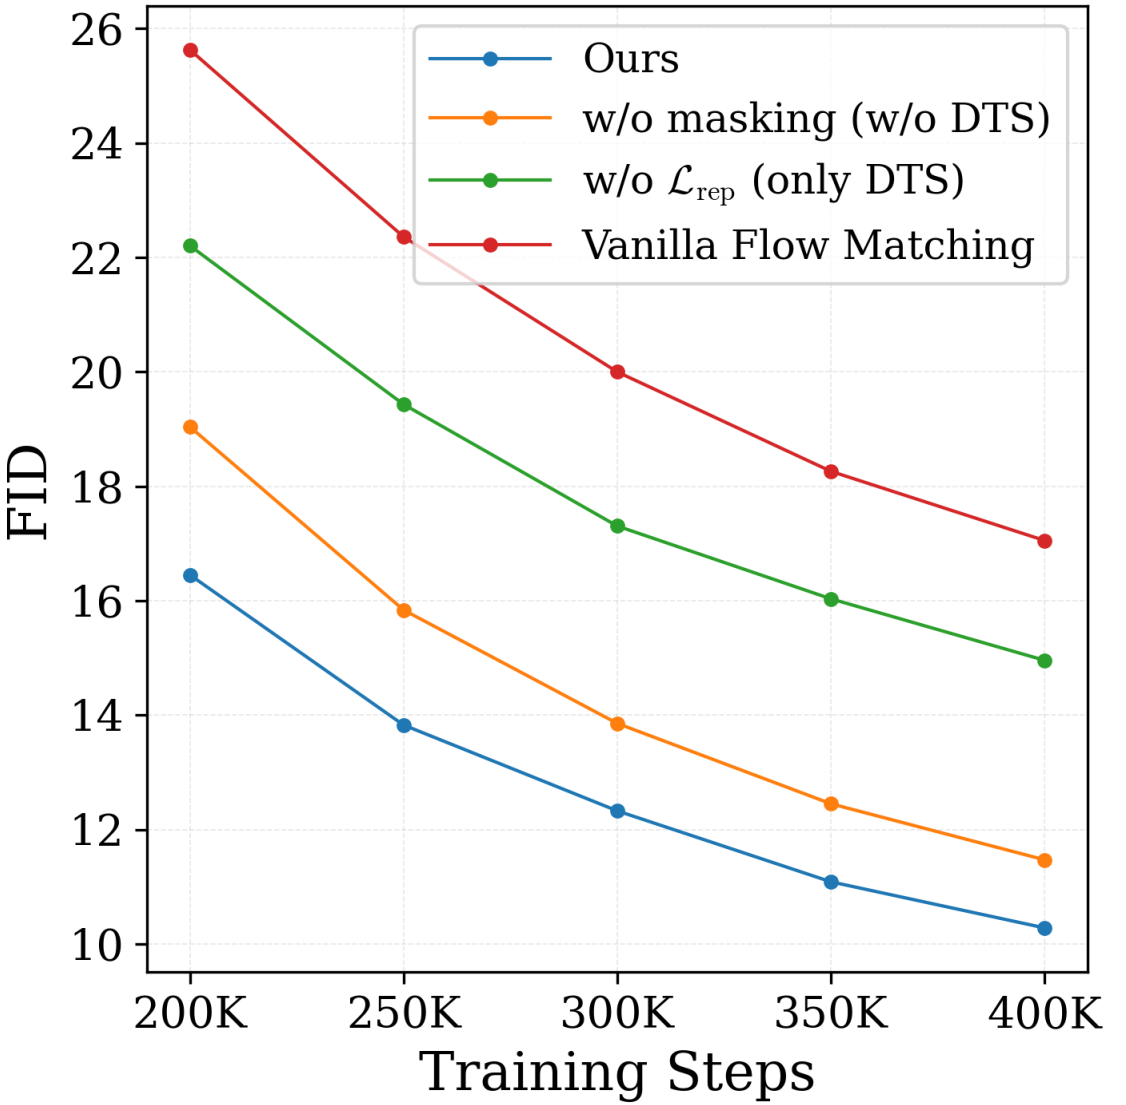

Ablation Study. Removing the self-supervised loss ($\mathcal{L}_{rep}$) is most detrimental. Removing Dual-Timestep Scheduling (w/o masking) significantly harms the results. Dual-Timestep Scheduling alone slightly improves the generation quality.